


Artificial Intelligence

Intelligent Agents

Dr. Bilgin Avenoğlu

Agents and Environments


- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and acting upon that environment through actuators.



Agents interact with environments through sensors and actuators.

Example - Vacuum Agent

- The vacuum agent perceives which square it is in and whether there is dirt in the square.
- The agent starts in square A.
- The available actions are to
 - move to the right,
 - move to the left,
 - suck up the dirt,
 - do nothing.
- One very simple agent function is the following: **if the current square is dirty, then suck; otherwise, move to the other square.**



Intelligent?

- What is the right way to fill out the table?
- In other words, what makes an agent **good** or **bad**, **intelligent** or **stupid**?

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>
<i>[A, Clean], [A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

Figure 2.3 Partial tabulation of a simple agent function for the vacuum-cleaner world shown in Figure 2.2. The agent cleans the current square if it is dirty, otherwise it moves to the other square. Note that the table is of unbounded size unless there is a restriction on the length of possible percept sequences.

The Concept of Rationality

- A rational agent is one that **does** the **right thing**.
- Obviously, doing the right thing is better than doing the wrong thing, but **what does it mean to do the right thing?**

Performance Measures

- When an agent is plunked down in an environment, it generates a **sequence of actions** according to the percepts it receives.
- This sequence of actions causes the environment to go through a **sequence of states**.
- If the **sequence is desirable**, then the agent has performed well.
- This notion of desirability is captured by a **performance measure** that evaluates any given sequence of environment states.
 - amount of dirt cleaned up in a single eight-hour shift.
 - a rational agent can maximize this performance measure by **cleaning** up the dirt, **then dumping** it all on the floor, then cleaning it up again, and so on.
 - reward the agent for having a clean floor.
 - one point could be awarded for each clean square at each time step (perhaps with a penalty for electricity consumed and noise generated).

Rationality

- What is rational at any given time depends on four things:
 - The **performance measure** that defines the criterion of success.
 - The agent's **prior knowledge** of the environment.
 - The **actions** that the agent can perform.
 - The agent's **percept sequence** to date.
- This leads to a definition of a rational agent:
 - For **each possible percept sequence**, a rational agent should **select an action** that is expected to **maximize its performance** measure, given the evidence provided by the percept sequence and whatever **built-in knowledge** the agent has.

Rationality

- The performance measure awards **one point** for each clean square at each time step, over a “lifetime” of 1000 time steps.
- The “geography” of the **environment** is **known a priori** but the dirt distribution and the initial location of the agent are not.
 - Clean squares stay clean and sucking cleans the current square. The Right and Left actions move the agent one square except when this would take the agent outside the environment, in which case the agent remains where it is.
- The only available actions are **Right**, **Left**, and **Suck**.
- The agent **correctly perceives** its **location** and whether that location contains **dirt**.
- Under these circumstances the agent is indeed rational; its expected performance is at least as good as any other agent's.

Omniscience, Learning, and Autonomy

- An **omniscient** agent knows the actual outcome of its actions and can act accordingly; but omniscience is **impossible** in reality.
 - information gathering—is an important part of rationality
 - exploration
- A rational agent should not only gather information but also to **learn** as much as possible from what it perceives.
- To the extent that an agent relies on the **prior knowledge** of its designer rather than on its **own percepts** and **learning processes**, we say that the agent lacks **autonomy**.

Example

- The female **sphex** will dig a **burrow**, go out and sting a **caterpillar** and **drag** it to the burrow, enter the burrow again to **check** all is well, drag the caterpillar inside, and lay its eggs.
- The caterpillar serves as a food source when the eggs hatch.
- If an entomologist **moves** the **caterpillar** a few inches away while the sphex is doing the check, it will **revert** to the “drag the caterpillar” step of its plan and **will continue the plan without modification, re-checking the burrow**, even after dozens of caterpillar-moving interventions.
- The sphex is **unable to learn** that its **innate plan is failing**, and thus will not change it.

Task Environment

- The **performance measure**, the **environment**, and the **agent's actuators** and **sensors**.
 - We group all these under the heading of the **task environment**.
 - **PEAS** (Performance, Environment, Actuators, Sensors)

PEAS - automated taxi

- To design a **rational agent**, we must **specify** the **task environment**
- Consider, e.g., the task of designing an automated taxi:
 - **Performance measure??** safety, destination, profits, legality, comfort, . . .
 - **Environment??** US streets/freeways, traffic, pedestrians, weather, . . .
 - **Actuators??** steering, accelerator, brake, horn, speaker/display, . . .
 - **Sensors??** video, accelerometers, gauges, engine sensors, keyboard, GPS, . . .

PEAS - Internet shopping agent

- **Performance measure??** price, quality, appropriateness, efficiency
- **Environment??** current and future WWW sites, vendors, shippers
- **Actuators??** display to user, follow URL, fill in form
- **Sensors??** HTML pages (text, graphics, scripts)

Properties of Task Environments

- **Fully observable** vs. **partially observable**
 - noisy and inaccurate sensors, unknown maps etc.
- **Single-agent** vs. **multiagent**
 - a crossword puzzle - single
 - playing chess - two-agents
- **Competitive** vs. **cooperative**
- **Deterministic** vs. **nondeterministic**
 - if the next state of the environment is completely determined - deterministic.
 - taxi driving - nondeterministic,
 - vacuum world - deterministic
- **Episodic** vs. **sequential**
 - the next episode does not depend on the actions taken in previous episodes
 - defective parts on an assembly - episodic
 - chess and taxi driving are sequential
- **Static** vs. **dynamic**
 - taxi driving is clearly dynamic
 - crossword puzzles are static
- **Discrete** vs. **continuous**
 - the speed and location of the taxi
- **Known** vs. **unknown**
 - in a known environment, the outcomes for all actions are given.

Environment Types

- The environment type largely determines the agent design
- The **real world** is (of course) **partially observable, stochastic, sequential, dynamic, continuous, multi-agent**

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??	No	No	No	No
Static??	Yes	Semi	Semi	No
Discrete??	Yes	Yes	Yes	No
Single-agent??	Yes	No	Yes (except auctions)	No

Agent Types

- Four basic types in order of increasing **generality**:
 - **simple reflex** agents
 - **reflex agents** with **state**
 - **goal-based** agents
 - **utility-based** agents
- All these can be turned into **learning agents**

Simple Reflex Agents

- These agents **select actions** on the basis of the **current percept**, ignoring the **rest** of the percept history.

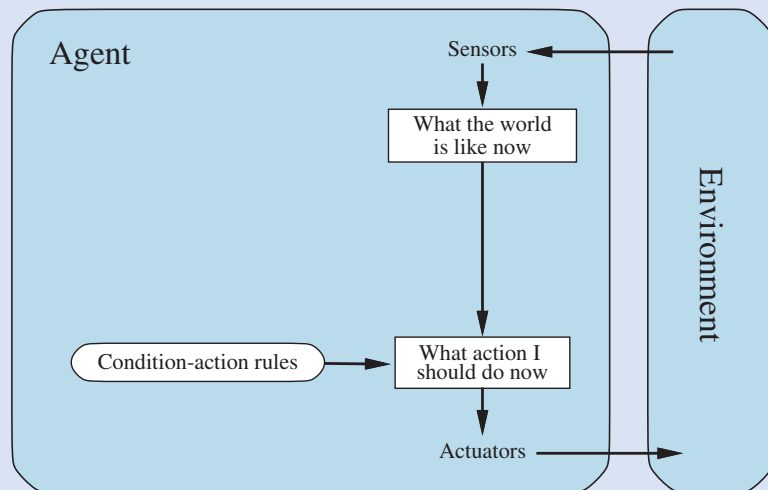


Figure 2.9 Schematic diagram of a simple reflex agent. We use rectangles to denote the current internal state of the agent's decision process, and ovals to represent the background information used in the process.

if *car-in-front-is-braking* **then** *initiate-braking*.

```
function Reflex-Vacuum-Agent([location, status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

function SIMPLE-REFLEX-AGENT(*percept*) **returns** an action

persistent: *rules*, a set of condition–action rules

state ← INTERPRET-INPUT(*percept*)

rule ← RULE-MATCH(*state*, *rules*)

action ← *rule*.ACTION

return *action*

Figure 2.10 A simple reflex agent. It acts according to a rule whose condition matches the current state, as defined by the percept.

Model-based Reflex Agents

- To handle **partial observability** - keep **track** of the part of the world that the agent can't see now.
- **Maintain** some sort of **internal state** that **depends** on the **percept history** and thereby reflects at least some of the unobserved aspects of the current state.

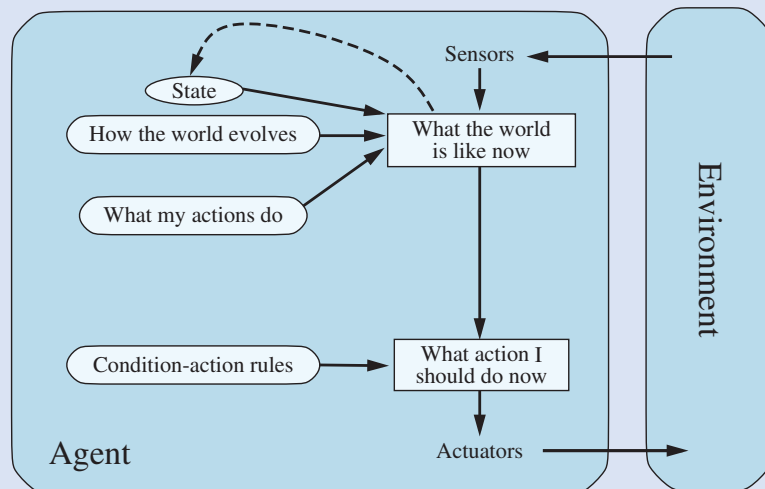


Figure 2.11 A model-based reflex agent.

function MODEL-BASED-REFLEX-AGENT(*percept*) **returns** an action

persistent: *state*, the agent's current conception of the world state
transition_model, a description of how the next state depends on the current state and action
sensor_model, a description of how the current world state is reflected in the agent's percepts
rules, a set of condition-action rules
action, the most recent action, initially none

state ← UPDATE-STATE(*state*, *action*, *percept*, *transition_model*, *sensor_model*)
rule ← RULE-MATCH(*state*, *rules*)
action ← *rule*.ACTION
return *action*

Figure 2.12 A model-based reflex agent. It keeps track of the current state of the world, using an internal model. It then chooses an action in the same way as the reflex agent.

Goal-based Agents

- As well as a **current state**, the agent needs some sort of **goal** information that describes **situations** that are **desirable**
 - for example, being at a particular destination.
- The agent program can **combine the goal** with the **model** to **choose actions** that achieve the goal.
- **Search** (Chapters 3, 4, and 6) and **planning** (Chapter 11) are the subfields of AI devoted to finding action sequences that achieve the agent's goals.

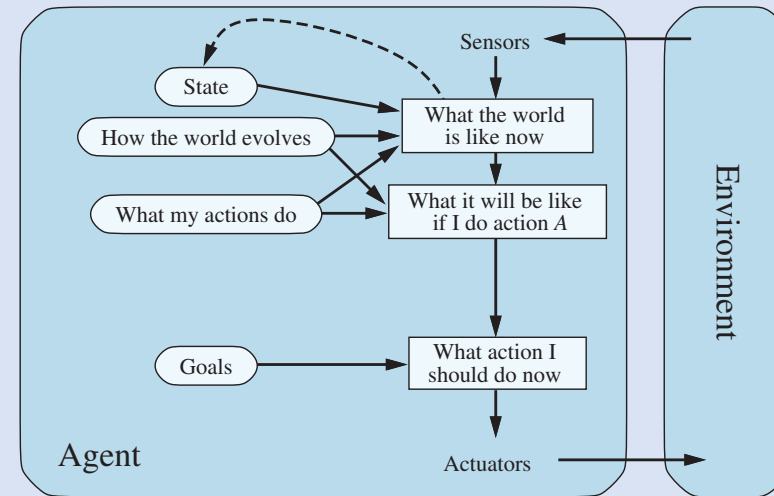


Figure 2.13 A model-based, goal-based agent. It keeps track of the world state as well as a set of goals it is trying to achieve, and chooses an action that will (eventually) lead to the achievement of its goals.

Utility-based Agents

- **Goals** alone are **not enough** to generate high-quality behavior in most environments.
 - For example, many action sequences will get the taxi to its destination, but some are quicker, safer, more reliable, or cheaper than others.

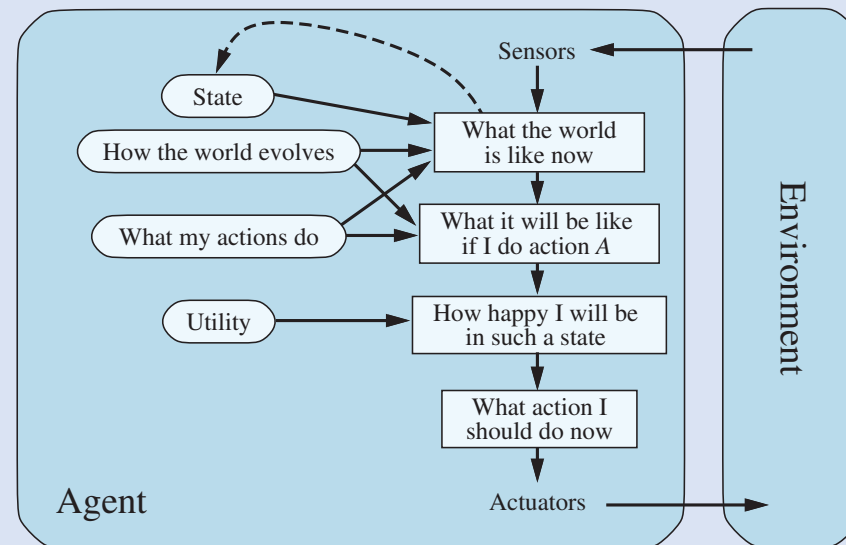


Figure 2.14 A model-based, utility-based agent. It uses a model of the world, along with a utility function that measures its preferences among states of the world. Then it chooses the action that leads to the best expected utility, where expected utility is computed by averaging over all possible outcome states, weighted by the probability of the outcome.

Learning Agents

- **Learning element** is responsible for making improvements.
- The **performance element**, or agent: takes **percepts** and decides on **actions**.
- The learning element
 - uses feedback from the **critic** on **how the agent is doing**
 - determines how the **performance element should be modified** to do better in the future.

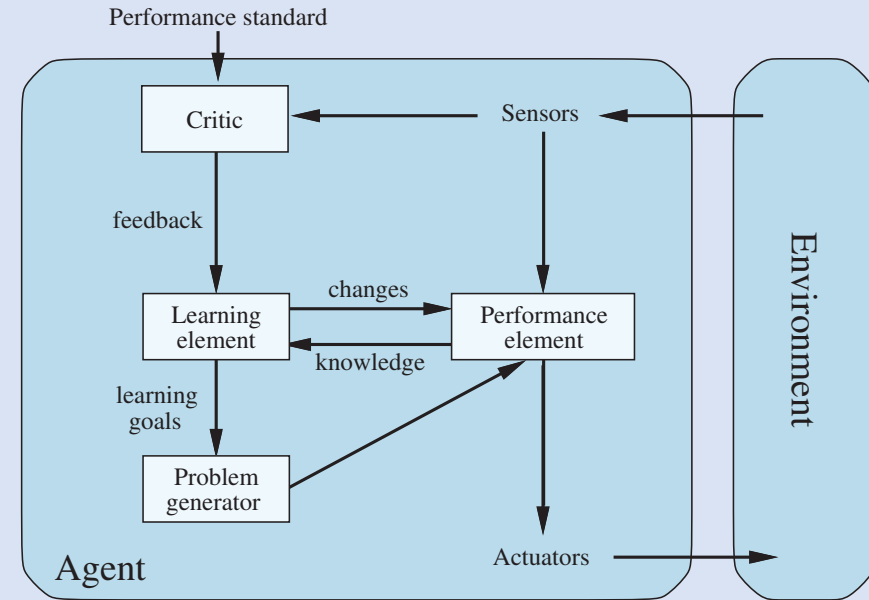


Figure 2.15 A general learning agent. The “performance element” box represents what we have previously considered to be the whole agent program. Now, the “learning element” box gets to modify that program to improve its performance.

- **Problem generator** is responsible for suggesting actions that will lead to new and informative experiences.

How the Components of Agent Programs Work?

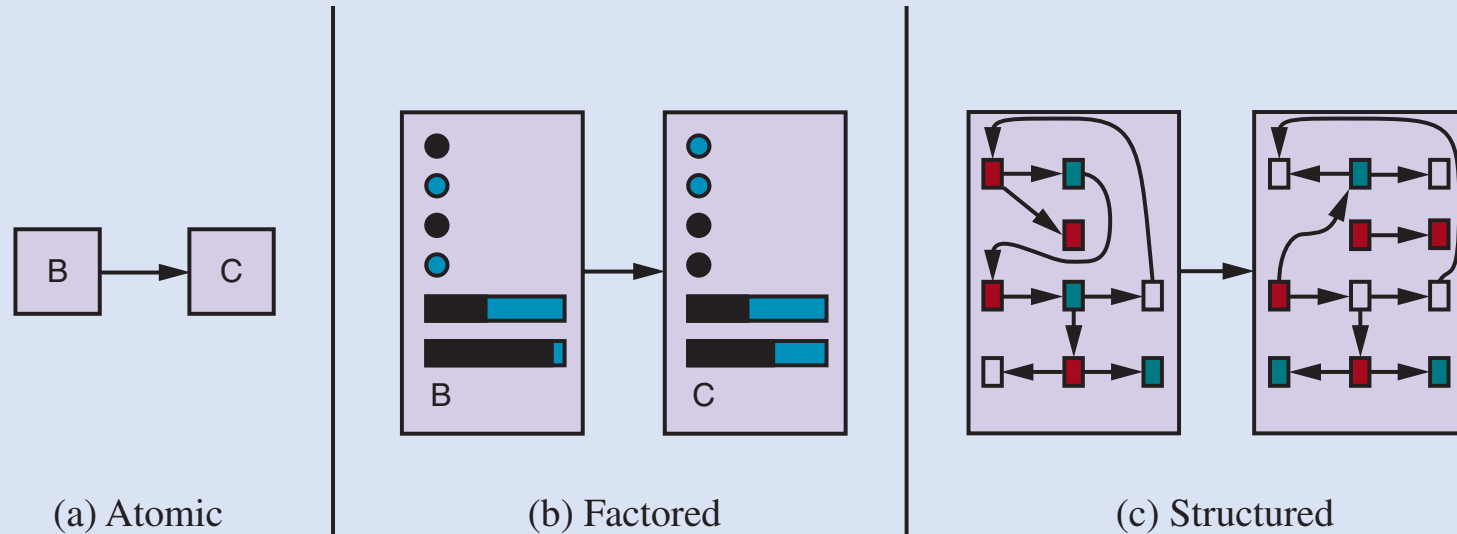



Figure 2.16 Three ways to represent states and the transitions between them. (a) Atomic representation: a state (such as B or C) is a black box with no internal structure; (b) Factored representation: a state consists of a vector of attribute values; values can be Boolean, real-valued, or one of a fixed set of symbols. (c) Structured representation: a state includes objects, each of which may have attributes of its own as well as relationships to other objects.

How the Components of Agent Programs Work?

- In an **atomic representation** each state of the world is **indivisible** - it has no internal structure.
 - Search and game-playing (Chapters 3, 4, and 6),
 - Hidden Markov models (Chapter 14),
 - Markov decision processes (Chapter 16).
- A **factored representation** splits up each state into a fixed set of **variables** or **attributes**, each of which can have a **value**.
 - Constraint satisfaction algorithms (Chapter 5),
 - Propositional logic (Chapter 7),
 - Planning (Chapter 11),
 - Bayesian networks (Chapters 12, 13, 14, 15, and 18),
 - Machine learning algorithms.
- In **structured representation**, **objects** and their various and varying **relationships** can be described explicitly.
 - relational databases and first-order logic (Chapters 8, 9, and 10),
 - first-order probability models (Chapter 18),
 - natural language understanding (Chapters 24 and 25).
 - In fact, much of what humans express in natural language concerns objects and their relationships.

The End!



[This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)